

# Squeezing Your Fine-Tuning Data to the Last Drop

## From Selection to Rebalancing

Minghao Wu

Department of Data Science and Artificial Intelligence  
Monash University

November 14, 2025

## About Me

- Final-year Ph.D. candidate at Monash University
- Research focus: Large Language Models, Multilinguality, Machine Translation
- 20+ papers in top-tier venues (ICML, ACL, EMNLP, COLING, TACL, etc.)
- Outstanding Paper Award at ACL 2025
- Visit/Intern experience: Huawei, Tencent, Alibaba, MBZUAI
- <https://minghao-wu.github.io/>

# Overview

1. Introduction
2. GraphFilter: Static Data Selection
3. Mixture-of-Skills: Dynamic Data Rebalancing
4. Conclusion

# The Data Challenge in LLM Fine-Tuning

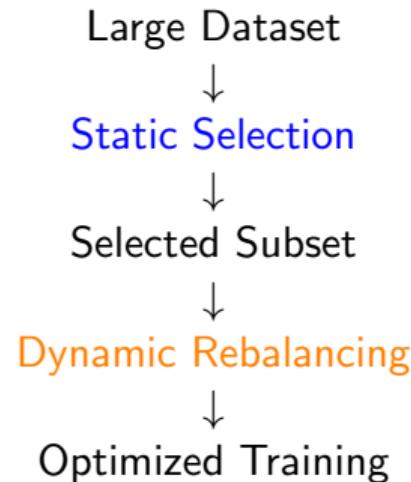
## Why does data matter in LLM fine-tuning?

- Training data directly shapes model capabilities
- Quality determines how well the model learns
- Diversity ensures comprehensive skill coverage
- Scale vs. efficiency trade-offs

## Key Challenges:

- Selection: What data to train on?
- Composition: How to mix datasets?
- Optimization: When to use what data?

## The Pipeline



*Two complementary approaches*

# Two Complementary Approaches

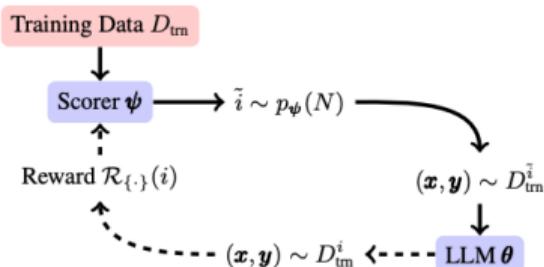
## Study 1: GraphFilter

- Static Data Selection
- Bipartite graph representation
- Set cover optimization
- Quality  $\times$  Diversity priority



## Study 2: Mixture-of-Skills

- Dynamic Rebalancing
- Reinforcement learning framework
- Adaptive data rebalancing
- Skill-aware training



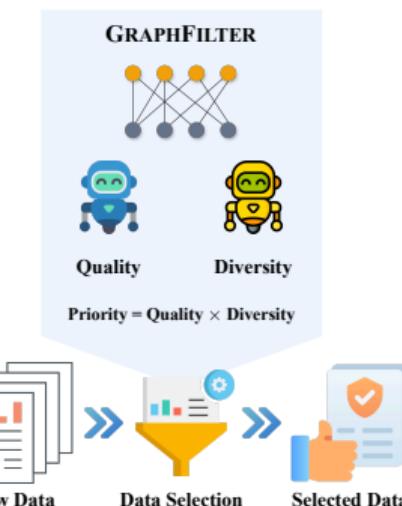
## The Best of Both Worlds: Bridging Quality and Diversity in Data Selection with Bipartite Graph

Minghao Wu, Thuy-Trang Vu, Lizhen Qu, Gholamreza Haffari

2025 Forty-Second International Conference on Machine Learning

### Key Contributions:

- Novel formulation as set cover problem
- Bipartite graph representation
- Multiplicative priority function
- Outperforms 9 baselines on 6 benchmarks



# The Quality-Diversity Dilemma

## Existing Methods Fall Short:

- Quality-focused: Select high-scoring examples
- Diversity-focused: Maximize coverage
- **Problem:** One aspect sacrificed for the other

## Real-world Analogy:

- Curating a library collection
- Want both high-quality books *and* diverse topics
- Balance is key for comprehensive learning

## The Challenge



**GraphFilter: Best of Both Worlds**

# What is the Set Cover Problem?

## Classic Computer Science Problem:

- Given a universe of elements  $U$
- Given a collection of sets  $\mathcal{S} = \{S_1, S_2, \dots, S_m\}$
- Each set  $S_i \subseteq U$
- **Goal:** Find minimum number of sets that cover all elements in  $U$

## NP-Hard Problem:

- No polynomial-time exact solution
- Greedy approximation works well
- Widely applicable in real-world scenarios

## Mathematical Formulation

$$\begin{aligned} & \min \sum_{i=1}^m x_i \\ & \text{subject to:} \\ & \sum_{i:e \in S_i} x_i \geq 1, \forall e \in U \\ & x_i \in \{0, 1\} \end{aligned}$$

where  $x_i = 1$  if set  $S_i$  is selected

# From Set Cover to Data Selection

## The Connection:

- Universe  $\mathcal{U}$ : All possible n-grams
- Sets  $\mathcal{S}$ : Training sentences
- Coverage: Each sentence covers its n-grams
- Goal: Select sentences that cover diverse n-grams

## Why This Makes Sense:

- N-grams represent linguistic patterns
- Diverse n-gram coverage = diverse training data
- Natural formulation for diversity

## Data Selection as Set Cover

**Sentence 1:** "How to cook pasta"

**N-grams:** {how, to, cook, pasta, how\_to, to\_cook, ...}

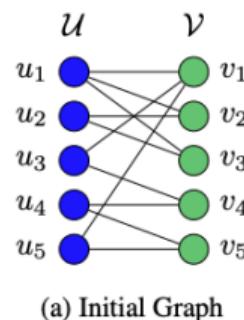
**Sentence 2:** "How to bake bread"

**N-grams:** {how, to, bake, bread, how\_to, to\_bake, ...}

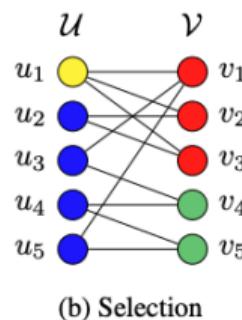
# Bipartite Graph Representation

## Graph Components:

- **Left nodes:** Sentences  $u_1, u_2, \dots, u_5$
- **Right nodes:** N-grams  $v_1, v_2, \dots, v_5$
- **Edges:** Connect sentences to their n-grams



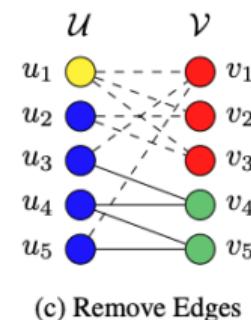
(a) Initial Graph



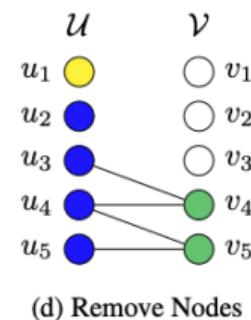
(b) Selection

## Selection Process:

- Select sentence with highest priority
- Remove covered n-grams from graph
- Update priorities dynamically
- Repeat until budget exhausted



(c) Remove Edges



(d) Remove Nodes

Figure: Example: Selecting  $u_1$  covers n-grams  $\{v_1, v_2, v_3\}$ , then selecting  $u_4$  covers remaining n-grams

But wait... This only considers diversity, not quality!

# Incorporating Quality: The Priority Function

## Priority Function Design:

$$\phi(u) = \text{Quality}(u) \times \max(1, \text{Diversity}(u))$$

## Quality Metric - IFD Score:

$$\text{Quality}(u) = \text{PPL}(y|x)/\text{PPL}(y)$$

## Diversity Metric - TF-IDF:

$$\text{Diversity}(u) = \sum_{v \in \mathcal{V}_u} \text{TF-IDF}(v)$$

## Why Multiplicative?

High Quality + Low Diversity  
⇒ Low Priority

Low Quality + High Diversity  
⇒ Low Priority

High Quality + High Diversity  
⇒ High Priority

*Both aspects must be strong for selection!*

# GraphFilter Algorithm in Action

## Iterative Selection Process:

1. **Initialize:** Empty selection  $S = \emptyset$
2. **Compute:** Priority  $\phi(u)$
3. **Select:**  $u^*$  with highest priority
4. **Update:** Remove  $u^*$  and n-grams
5. **Repeat:** Until budget  $k$  is reached

## Dynamic Priority Updates:

- Priorities change as n-grams are covered
- Encourages selection of complementary sentences
- Balances quality and remaining diversity

## Computational Efficiency

- Max-heap  $\mathcal{O}(\log N)$  per iteration
- Localized priority updates
- Scalable to large datasets

## N-gram Selection

- Use n-grams up to length 3
- Balances granularity and efficiency
- Captures meaningful linguistic patterns
- Empirically effective

# Experimental Setup

## Training Dataset:

- Magpie dataset (300K instances)
- High-quality instruction-response pairs
- Select 10K subset

## Model Backbones:

- Gemma-2-2B
- Mistral-7B-v0.3
- Llama-3-8B

## Baseline Methods (9 total):

- Heuristic: Random, Longest
- Quality-based: PPL, ArmoRM, AlpaGagus, DEITA, IFD
- Diversity-based: K-means, InsTag

## Evaluation Benchmarks (6 total):

- MMLU, ARC, HellaSwag, GSM8K
- AlpacaEval-2.0, MT-Bench using GPT-4o

# Main Results: GraphFilter Outperforms All Baselines

## Consistent Improvements Against Random:

- **Gemma-2-2B**: Up to **+2.03**
- **Mistral-7B**: Up to **+2.83**
- **Llama-3-8B**: Up to **+2.46**

## Key Findings:

- GraphFilter achieves best/second-best on most benchmarks
- Quality-only methods show **benchmark bias** (e.g., ArmoRM good on AlpacaEval, poor elsewhere)

	Standard	LLM	ALL
Random	47.75	41.04	45.51
Longest	46.91	39.96	44.59
Perplexity	48.27	40.28	45.61
ArmoRM	48.21	<u>42.66</u>	46.36
AlpaGasus	48.96	41.90	46.60
DEITA	48.78	41.70	46.42
SuperFilter	49.10	41.91	46.70
Kmeans	48.90	41.72	46.51
InsTag	<u>49.93</u>	41.72	<u>47.19</u>
GraphFilter (Ours)	<b>50.55</b>	<b>42.79</b>	<b>47.97</b>

Table: Main results given by Llama-3-8B on the standardized benchmarks and LLM-as-a-Judge benchmarks. The best results are highlighted in **bold**, and the second-best results are highlighted in underline.

# Main Results: GraphFilter is Computationally Efficient

## Computational Efficiency

- Max-heap  $\mathcal{O}(\log N)$  per iteration
- Localized priority updates
- Scalable to large datasets

	Runtime (hrs)
Perplexity	0.92
ArmoRM	5.93
AlpaGasus	32.34
DEITA	22.65
SuperFilter	1.95
Kmeans	2.26
InsTag	25.48
GraphFilter (Ours)	2.48
w/o priority $\phi(u)$	0.53 <sup>†</sup>

Table: Runtime (in hours) for selecting 10K training instances. <sup>†</sup> indicate the CPU-only method.

# Analysis: GraphFilter Balances Quality and Diversity

## Visualization Analysis:

- **Lexical Diversity:** Measured by MTLD metric
- **Data Quality:** Assessed by SkyworkRM

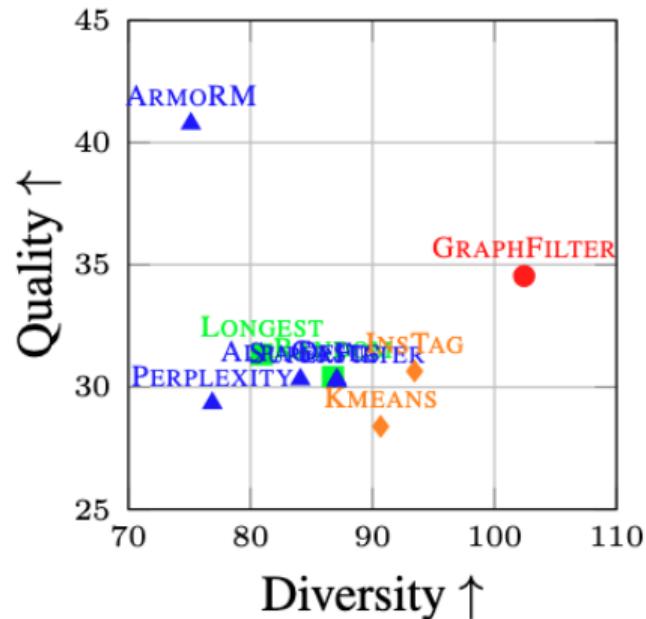
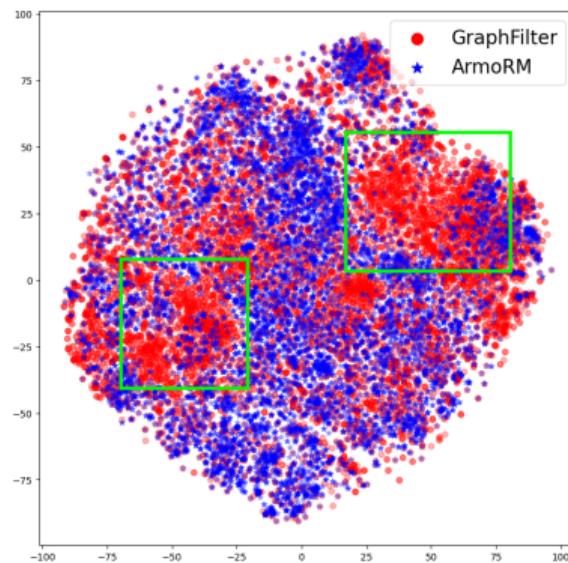


Figure: GraphFilter achieves **highest lexical diversity** and **second-best data quality** among all methods.

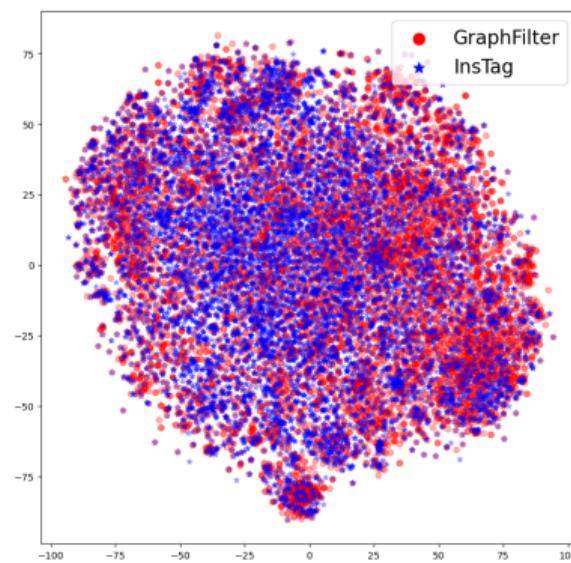
# Analysis: GraphFilter Balances Quality and Diversity

## Semantic Visualization: t-SNE plots using BGE embeddings

- Selects instructions not chosen by quality-only methods
- Similar semantic diversity to diversity-focused methods



(a) GraphFilter vs. ArmoRM (Quality only)



(b) GraphFilter vs. InsTag (Diversity only)

# Analysis: N-gram Combinations Capture Multi-Level Features

## Different N-grams Capture Different Patterns:

- **Unigrams (n=1)**: Individual words, basic vocabulary
- **Bigrams (n=2)**: Local word relationships
- **Trigrams (n=3)**: Phrasal patterns, syntax

## Key Finding:

- Combining 1-grams + 2-grams + 3-grams significantly outperforms individual n-gram types
- Each level provides complementary information
- Integration consolidates features effectively

## Performance by N-gram Type

N-gram	ALL
1-gram only	46.48
2-gram only	46.63
3-gram only	47.15
1+2+3-grams	<b>47.97</b>

Multi-level features are essential!

# Analysis: Trigrams Provide Optimal Balance

## N-gram Size Analysis ( $n_{max}=1$ to 5):

- **Performance:** Peaks at trigrams ( $n_{max}=3$ )
- **Efficiency:** Decreases with larger  $n_{max}$
- **Diminishing returns:** Beyond  $n_{max}=3$

## Why Trigrams Work Best:

- Capture meaningful linguistic patterns
- Balance granularity and computational cost
- Avoid over-specification of larger n-grams

## N-gram Size Trade-offs

$n_{max}$	ALL	Time (hrs)
1	46.48	2.12
2	47.31	2.30
<b>3</b>	<b>47.97</b>	<b>2.48</b>
4	47.43	3.38
5	47.85	4.58

Sweet spot at  $n=3$

# N-grams Nodes:  
0.1M (1-gram) → 7.4M (5-gram)

# Analysis: Instruction Diversity Matters Most

## What to Apply GraphFilter To?

- Each training instance has **instruction** + **response**
- Tested three scenarios:
  - Instructions only
  - Responses only
  - Both instructions and responses

## Surprising Result:

- **Instructions only** performs best
- Instruction diversity more impactful than response diversity
- Quality remains similar
- **Rethinking**: Response diversity also matters!

Type	Benchmarks			Lexical Diversity		Quality
	Standard	LLM	ALL	Instruction	Response	
Instruction	<b>50.55</b>	<b>42.79</b>	<b>47.97</b>	<b>102.43</b>	71.74	<b>81.54</b>
Response	47.16	39.71	44.68	90.22	<b>73.57</b>	81.52
Inst. + Resp.	48.03	41.20	45.76	90.13	72.60	81.52

# Analysis: Budget Determines Quality vs. Diversity Priority

## When to Prioritize What?

- **Small budgets (1K, 5K):** Quality methods (SuperFilter) excel
- **Large budgets (10K+):** Diversity methods (InsTag) catch up
- **All budgets:** GraphFilter consistently demonstrates performance gains

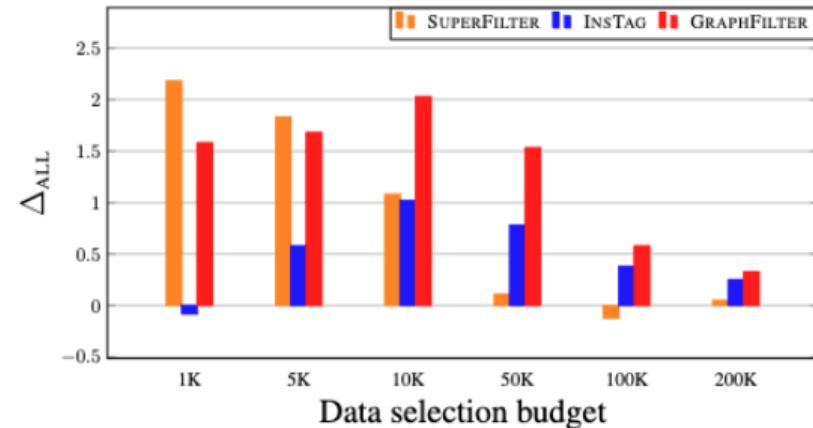


Figure: Relative gains against Random. Budget size affects strategy effectiveness

# From Static Selection to Dynamic Optimization

## GraphFilter Insights:

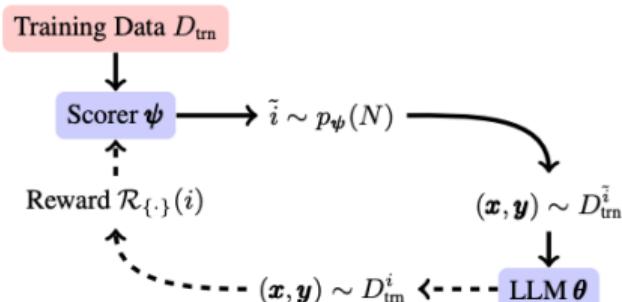
- Quality  $\times$  Diversity balance is crucial
- Static selection has limitations
- What about during training?

## New Challenge:

- Models learn different skills at different rates
- Static data mixing may not be optimal
- Finding the right data mix is expensive

## The Next Question

Can we optimize data usage *dynamically* during fine-tuning?



Enter: Mixture-of-Skills

# Study 2: Mixture-of-Skills

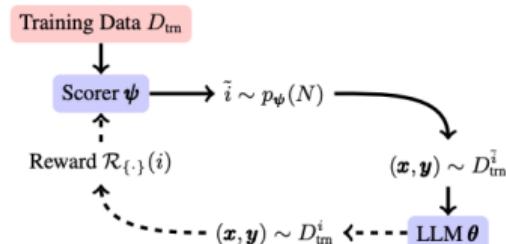
## Mixture-of-Skills: Learning to Optimize Data Usage for Fine-Tuning Large Language Models

Minghao Wu, Thuy-Trang Vu, Lizhen Qu, Gholamreza Haffari

2024 Conference on Empirical Methods in Natural Language Processing

### Key Contributions:

- Reinforcement learning framework for data rebalancing
- Model-agnostic dynamic optimization
- Adaptive skill development during training
- MoSpec extension for task-specific fine-tuning



# The Data Mixing Challenge

## LLMs Need Multiple Skills:

- Writing, coding, mathematics, chatting, etc.
- Each skill requires different training data
- Datasets are **heterogeneous** and **imbalanced**

## Current Problems:

- **Static mixing:** Ignores learning dynamics
- **Data capping:** Limits large dataset utilization
- **Cost:** Finding optimal ratios is costly

## Example Challenge

- Math: 10K samples
- Code: 100K samples
- Chat: 1M samples

## Learning Dynamics:

- Some skills learned quickly
- Others need more exposure
- Skills can interfere

# Mixture-of-Skills: Bilevel Optimization

## The Framework:

- Outer level: Optimize LLM parameters  $\theta$
- Inner level: Optimize data sampling via scorer network  $\psi$
- Goal: Learn optimal data usage automatically

## Mathematical Formulation:

$$\psi = \underset{\psi}{\operatorname{argmin}} \mathcal{J}(D_{\text{trn}}; \theta(\psi))$$

$$\theta(\psi) = \underset{\theta}{\operatorname{argmin}} \mathbb{E}_{i \sim p_{\psi}(N)} [\mathcal{L}(D_{\text{trn}}^i; \theta)]$$

## Why RL?

- Bilevel optimization not directly differentiable
- REINFORCE algorithm handles discrete sampling
- Enables dynamic adaptation

# MoS Algorithm: Step-by-Step Walkthrough

## Algorithm Overview:

1. **Initialize:** Start with uniform sampling ( $\tau = \infty$ )
2. **Training Loop:** For each step  $t$ :
  - Sample dataset  $\tilde{i} \sim p_\psi(N)$
  - Sample batch from  $D_{\text{trn}}^{\tilde{i}}$
  - Update LLM:  $\theta \leftarrow \theta - \alpha \nabla_\theta \mathcal{L}$
3. **Scorer Update** (every  $S$  steps):
  - Sample batch from each dataset
  - Compute rewards  $\mathcal{R}(i)$  for each dataset
  - Update scorer:  
$$\psi \leftarrow \psi + \gamma \sum_i \mathcal{R}(i) \nabla_\psi \log p_\psi(i)$$

## Key Components:

- **Scorer Network  $\psi$ :** Simple 2-layer MLP
- **Update Frequency  $S$ :** Computational efficiency
- **Rewards  $\mathcal{R}(i)$ :** Guide optimization

## Dynamic Adaptation:

- Sampling probabilities change over time
- Responds to current model state
- Balances different skills automatically

**But how do we design rewards?**

# Reward Design: Three Perspectives

## 1. Transferability ( $\mathcal{R}_{\text{cosine}}$ ):

- Measure similarity between datasets
- Use mini-batch embeddings from LLM
- Higher similarity = better transferability

$$\mathcal{R}_{\text{cosine}}(i) = \frac{1}{N} \sum_{n=1}^N \frac{\mathbf{z}^i \cdot \mathbf{z}^n}{\|\mathbf{z}^i\| \cdot \|\mathbf{z}^n\|}$$

## 2. Difficulty ( $\mathcal{R}_{\text{diff}}$ ):

- Relative perplexity decrease after fine-tuning
- Higher values = more difficult datasets need more attention

$$\mathcal{R}_{\text{diff}}(i) = \frac{1}{L} \sum_{j=1}^L \frac{\text{PPL}(\mathbf{y}_j^i; \mathbf{x}_j^i, \boldsymbol{\theta})}{\text{PPL}(\mathbf{y}_j^i; \mathbf{x}_j^i, \boldsymbol{\theta}_0)}$$

## 3. Learning Trajectory:

- Exponential Moving Average (EMA) for stability
- Smooths reward signals over time
- Prevents oscillations

$$\mathcal{R}(i) = \beta \mathcal{R}'(i) + (1 - \beta) \mathcal{R}''(i)$$

## Design Principles:

- **Transferability**: Knowledge sharing
- **Difficulty**: Learning progress
- **Trajectory**: Stability

*Rewards guide the scorer network to make informed sampling decisions*

# Experimental Setup

## Datasets (4 Skills):

- **Mathematics:** MathInstruct (comprehensive collection)
- **Medicine:** MedInstruct (medical instructions)
- **General:** ShareGPT (diverse, high-quality conversations)
- **NLP:** P3 (660 subsets, diverse NLP tasks)

## Model Backbones (3 total):

- Qwen1.5-0.5B
- Gemma-2B
- Llama-3-8B

## Baselines:

- **Heuristic:** Temperature sampling

$$q_\tau(i) = \frac{q(i)^{1/\tau}}{\sum_{n=1}^N q(n)^{1/\tau}}$$

- Proportional ( $\tau = 1$ ), Temperature ( $\tau = 10$ ), Uniform ( $\tau = \infty$ )

- **Dynamic:** MultiDDS, MultiUAT

## Evaluation:

- **MMLU:** 57 subjects (Math, Medicine, Others)
- **MT-Bench:** 8 skills (Coding, Writing, etc.)

# Main Results: MoS Consistently Outperforms Baselines

## Consistent Improvements:

- **Qwen1.5-0.5B**: +0.96 improvement
- **Gemma-2B**: +1.15 improvement
- **Llama-3-8B**: +2.45 improvement

## Key Findings:

- Larger models benefit more from MoS
- No universally optimal temperature  $\tau$
- Different rewards work better for different models
- EMA consistently helps performance

	MMLU	MT	ALL
Prop. ( $\tau = 1$ )	56.78	6.52	60.97
Temp. ( $\tau = 10$ )	56.17	6.66	61.40
Uni. ( $\tau = \infty$ )	55.72	6.63	60.99
MultiDDS	56.65	6.69	61.77
MultiUAT	55.66	6.67	61.18
MoS + cos	56.95	6.80	62.49
MoS + cos + EMA	<u>58.08</u>	<b>6.96</b>	<b>63.85</b>
MoS + diff	57.93	6.81	63.00
MoS + diff + EMA	<b>58.34</b>	<u>6.82</u>	<u>63.26</u>

Table: Results on MMLU and MT-Bench given by Llama-3-8B. Best results in **bold**, second-best underlined.

# Analysis: MoS Adapts Data Usage Dynamically

## Dynamic Sampling Behavior:

- Sampling probabilities evolve over training
- Reflects changing model needs
- Balances skill development

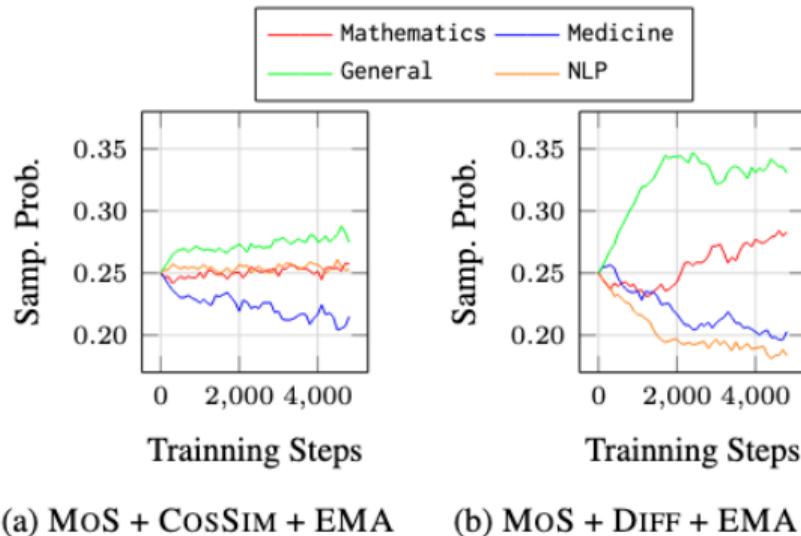


Figure: Sampling probabilities for each dataset over training steps. MoS adjusts focus dynamically.

# Analysis: MoS is Robust to Sampling Priors

## Initialization Sensitivity:

- MoS starts with uniform sampling ( $\tau = \infty$ )
- What if we start with different priors?
- **Key Finding:** MoS consistently outperforms best heuristic baseline regardless of initialization

## Robustness Analysis:

- All MoS variants beat Temperature ( $\tau = 10$ ) baseline
- Proper prior selection can further enhance performance
- **MoS is robust and effective**

	Prior $\tau$	MMLU	MT	ALL
Heuristic	10	56.17	6.66	61.40
MultiDDS	1	56.65	6.69	61.77
	10	56.88	6.58	61.34
	$\infty$	55.66	6.67	61.18
MoS	1	56.51	6.81	62.29
	10	58.22	<b>6.91</b>	<b>63.66</b>
	$\infty$	<b>58.34</b>	6.82	63.26

Table: MoS performance with different sampling priors using Llama-3-8B. All variants outperform the best heuristic baseline.

**Takeaway:** MoS learns to adapt regardless of how you start!

# Analysis: MoS is Compatible with Data Selection Methods

## Combining MoS with Instance Selection:

- Use Instruction-Following Difficulty (IFD) to select top 10% instances
- Apply MoS dynamic rebalancing on selected data
- **Best of both worlds:** Quality selection + dynamic optimization

## Complementary Effects:

- Static selection improves data quality
- MoS optimizes usage of selected data
- **Complementary approaches** enhance final performance

	MMLU	MT	ALL
Random			
Temp. ( $\tau = 10$ )	54.89	6.63	60.62
MoS	55.21	6.72	61.15
IFD Selection			
Temp. ( $\tau = 10$ )	55.13	6.69	61.02
MoS	<b>56.43</b>	<b>6.77</b>	<b>62.05</b>

Table: Combining MoS with instance selection (IFD) further improves performance using Llama-3-8B.

**Key Insight:** Static selection and dynamic rebalancing address different aspects of data optimization!

# MoSpec: Fine-tuning from Generalist to Specialist

## The Specialization Challenge:

- Large general-purpose models are costly to deploy
- Many applications need only narrow functionalities
- Smaller specialized models often outperform larger generalist ones

## MoSpec Approach:

- Harnesses diverse datasets to enhance target task performance
- Assigns higher rewards to target domain (e.g., mathematics)
- Learns optimal dataset distribution for specific capabilities

## Example: Math Specialization

- Target: Mathematics dataset
- Supporting: General, Medicine, NLP datasets
- **MoSpec + cosine**: Compute similarity between Math and other datasets
- **MoSpec + diff**: Double reward for Math dataset

**Key Insight:** SFT datasets from other domains are beneficial for the target task, especially when target dataset is incomplete

# MoSpec Results: Math Specialization Performance

## Experimental Setup:

- **Target:** Math specialization
- **Baselines:**
  - Temp. ( $\tau = 10$ )
  - MathLlama (trained only on Math)
- **Evaluation:** GSM8K, MATH, MMLU-math (M-math) benchmarks

## Key Findings:

- **MathLlama performs worst**
- **Other datasets are beneficial**
- **MoSpec learns optimal mixing** for specialization

	GSM8K	MATH	M-math	ALL
<i>Generalist</i>				
Temp. ( $\tau = 10$ )	49.62	9.54	28.36	29.17
MoS + cos + EMA	50.40	9.78	27.60	29.26
MoS + diff + EMA	49.58	10.26	<b>32.81</b>	30.88
<i>Math Specialist</i>				
MathLlama	41.02	9.76	30.34	27.04
MoSpec + cos + EMA	51.10	10.64	30.16	30.63
MoSpec + diff + EMA	<b>52.10</b>	<b>11.40</b>	32.24	<b>31.91</b>

Table: Math specialization results. MoSpec outperforms both generalist models and math-only training.

+4.87 improvement over math-only training!

# MoSpec Results: Sampling Behavior Analysis

## Dynamic Data Usage:

- MoSpec adjusts dataset sampling over training
- Increases focus on target domain (Math)
- Leverages supporting datasets effectively

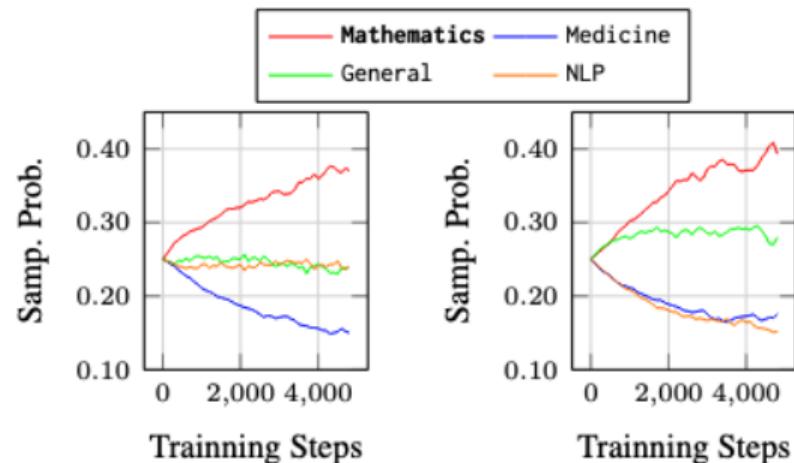


Figure: Sampling probabilities during Math specialization. **Left:** MoSpec + cos + EMA; **Right:** MoSpec + diff + EMA. Both increase focus on Math over time.

# Summary of Contributions

## GraphFilter: Principled Static Selection

- Novel set cover formulation for data selection
- Bipartite graph representation balancing quality × diversity
- Consistent improvements across models and benchmarks
- Computationally efficient and scalable

## MoS: Intelligent Dynamic Optimization

- Model-agnostic RL framework for data rebalancing
- Adaptive skill development during fine-tuning
- Substantial performance improvements
- MoSpec extension for task specialization

**Comprehensive solution covering the entire data optimization pipeline**

Thank You

## Questions & Discussion

### Contact Information:

- Email: [minghao.wu@monash.edu](mailto:minghao.wu@monash.edu)
- Twitter/X: [https://x.com/WuMinghao\\_nlp](https://x.com/WuMinghao_nlp)
- LinkedIn: <https://www.linkedin.com/in/minghao-wu-9087076a>
- GitHub: <https://github.com/minghao-wu>
- Website: <https://minghao-wu.github.io>
- Papers: Available on arXiv and conference proceedings

Special thanks to ILCC for the invitation and the opportunity to share this work